

## ml\_final\_project'

Ruchen Cai

3/16/2022

## import and clean the data

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.3.1 --
```

```
## v ggplot2 3.3.5      v purrr 0.3.4
```

```
## v tibble 3.1.2      v dplyr 1.0.6
```

```
## v tidyr      1.1.3      v stringr 1.4.0
```

```
## v readr    1.4.0    v forcats 0.5.1
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::filter() masks stats::filter()
```

```
## x dplyr::lag()      masks stats::lag()
```

```
load(file='nhanes2003-2004.rda')
```

```
# remove participants younger than 50 yo
```

```
NHANES.data <- subset(nhanes2003_2004, as.numeric(nhanes2003_2004$RIDAGEEX) > (50*12-1))
```

```
# make a matrix that contains only important predictors, excluding participants with missing data
variables <- c("RIDAGEYR", "RIAGENDR", "BPQ010", "BPQ060", "DIQ010", "DIQ050", "DIQ090", "MCQ010", "MCQ
```

```
NHANES.data <- na.omit(NHANES.data[variables])
```

```
# change all values to type numeric
```

```
for(i in 1:length(NHANES.data[1,])){
```

```
NHANES.data[,i] <- as.numeric(NHANES.data[,i])
```

}

```
summary(NHANES.data)
```

| ## | RIDAGEYR      | RIAGENDR      | BPQ010        | BPQ060        |
|----|---------------|---------------|---------------|---------------|
| ## | Min. :48.00   | Min. :1.000   | Min. :1.000   | Min. :1.000   |
| ## | 1st Qu.:58.00 | 1st Qu.:1.000 | 1st Qu.:1.000 | 1st Qu.:1.000 |
| ## | Median :64.00 | Median :1.000 | Median :1.000 | Median :1.000 |
| ## | Mean :64.25   | Mean :1.494   | Mean :1.267   | Mean :1.136   |
| ## | 3rd Qu.:72.00 | 3rd Qu.:2.000 | 3rd Qu.:1.000 | 3rd Qu.:1.000 |
| ## | Max. :83.00   | Max. :2.000   | Max. :5.000   | Max. :3.000   |
| ## | DIQ010        | DIQ050        | DIQ090        | MCQ010        |
| ## | Min. :1.000   | Min. :1.000   | Min. :1.000   | Min. :1.000   |
| ## | 1st Qu.:2.000 | 1st Qu.:2.000 | 1st Qu.:2.000 | 1st Qu.:2.000 |
| ## | Median :2.000 | Median :2.000 | Median :2.000 | Median :2.000 |
| ## | Mean :1.826   | Mean :1.956   | Mean :1.961   | Mean :1.899   |

|    |                |                |               |                |
|----|----------------|----------------|---------------|----------------|
| ## | 3rd Qu.:2.000  | 3rd Qu.:2.000  | 3rd Qu.:2.000 | 3rd Qu.:2.000  |
| ## | Max. :3.000    | Max. :2.000    | Max. :3.000   | Max. :3.000    |
| ## | MCQ053         | MCQ160A        | MCQ160B       | MCQ160K        |
| ## | Min. :1.000    | Min. :1.000    | Min. :1.000   | Min. :1.000    |
| ## | 1st Qu.:2.000  | 1st Qu.:1.000  | 1st Qu.:2.000 | 1st Qu.:2.000  |
| ## | Median :2.000  | Median :2.000  | Median :2.000 | Median :2.000  |
| ## | Mean :1.968    | Mean :1.547    | Mean :1.952   | Mean :1.921    |
| ## | 3rd Qu.:2.000  | 3rd Qu.:2.000  | 3rd Qu.:2.000 | 3rd Qu.:2.000  |
| ## | Max. :4.000    | Max. :3.000    | Max. :3.000   | Max. :3.000    |
| ## | MCQ160L        | BMXWAIST       | MCQ160M       | MCQ220         |
| ## | Min. :1.000    | Min. : 1.00    | Min. :1.000   | Min. :1.000    |
| ## | 1st Qu.:2.000  | 1st Qu.: 93.25 | 1st Qu.:2.000 | 1st Qu.:2.000  |
| ## | Median :2.000  | Median :308.00 | Median :2.000 | Median :2.000  |
| ## | Mean :1.955    | Mean :478.11   | Mean :1.861   | Mean :1.848    |
| ## | 3rd Qu.:2.000  | 3rd Qu.:895.00 | 3rd Qu.:2.000 | 3rd Qu.:2.000  |
| ## | Max. :3.000    | Max. :981.00   | Max. :3.000   | Max. :3.000    |
| ## | MCQ245A        | MCQ250A        | MCQ250B       | MCQ250C        |
| ## | Min. :1.000    | Min. :1.000    | Min. :1.000   | Min. :1.000    |
| ## | 1st Qu.:1.000  | 1st Qu.:1.000  | 1st Qu.:2.000 | 1st Qu.:2.000  |
| ## | Median :2.000  | Median :2.000  | Median :2.000 | Median :2.000  |
| ## | Mean :1.615    | Mean :1.523    | Mean :1.869   | Mean :1.836    |
| ## | 3rd Qu.:2.000  | 3rd Qu.:2.000  | 3rd Qu.:2.000 | 3rd Qu.:2.000  |
| ## | Max. :2.000    | Max. :3.000    | Max. :3.000   | Max. :3.000    |
| ## | MCQ250E        | MCQ250F        | MCQ250G       | MCQ265         |
| ## | Min. :1.000    | Min. :1.000    | Min. :1.000   | Min. :1.000    |
| ## | 1st Qu.:2.000  | 1st Qu.:2.000  | 1st Qu.:2.000 | 1st Qu.:2.000  |
| ## | Median :2.000  | Median :2.000  | Median :2.000 | Median :2.000  |
| ## | Mean :1.877    | Mean :1.804    | Mean :1.903   | Mean :1.908    |
| ## | 3rd Qu.:2.000  | 3rd Qu.:2.000  | 3rd Qu.:2.000 | 3rd Qu.:2.000  |
| ## | Max. :3.000    | Max. :3.000    | Max. :3.000   | Max. :4.000    |
| ## | SSQ011         | SSQ051         | WHQ030        | WHQ040         |
| ## | Min. :1.000    | Min. :1.000    | Min. :1.000   | Min. :1.000    |
| ## | 1st Qu.:1.000  | 1st Qu.:1.000  | 1st Qu.:1.000 | 1st Qu.:2.000  |
| ## | Median :1.000  | Median :1.000  | Median :1.000 | Median :2.000  |
| ## | Mean :1.093    | Mean :1.307    | Mean :1.779   | Mean :2.301    |
| ## | 3rd Qu.:1.000  | 3rd Qu.:2.000  | 3rd Qu.:3.000 | 3rd Qu.:3.000  |
| ## | Max. :4.000    | Max. :4.000    | Max. :5.000   | Max. :5.000    |
| ## | LBXRDW         | HSD010         | BPXPULS       | BPXML1         |
| ## | Min. : 5.00    | Min. :1.000    | Min. :1.000   | Min. : 2.000   |
| ## | 1st Qu.:18.00  | 1st Qu.:2.000  | 1st Qu.:1.000 | 1st Qu.: 5.000 |
| ## | Median :22.00  | Median :3.000  | Median :1.000 | Median : 6.000 |
| ## | Mean :24.55    | Mean :2.923    | Mean :1.085   | Mean : 6.559   |
| ## | 3rd Qu.:28.00  | 3rd Qu.:4.000  | 3rd Qu.:1.000 | 3rd Qu.: 8.000 |
| ## | Max. :97.00    | Max. :5.000    | Max. :2.000   | Max. :14.000   |
| ## | VIQ200         | BMXBMI         | BPXSY1        | BPXDI1         |
| ## | Min. :1.000    | Min. : 119     | Min. : 1.00   | Min. : 1.00    |
| ## | 1st Qu.:2.000  | 1st Qu.:1135   | 1st Qu.:12.00 | 1st Qu.:37.00  |
| ## | Median :2.000  | Median :1424   | Median :18.00 | Median :40.00  |
| ## | Mean :1.848    | Mean :1473     | Mean :20.64   | Mean :39.71    |
| ## | 3rd Qu.:2.000  | 3rd Qu.:1800   | 3rd Qu.:26.00 | 3rd Qu.:44.00  |
| ## | Max. :3.000    | Max. :2605     | Max. :77.00   | Max. :55.00    |
| ## | mortstat       |                |               |                |
| ## | Min. :0.0000   |                |               |                |
| ## | 1st Qu.:0.0000 |                |               |                |

```
## Median :0.0000
## Mean   :0.1796
## 3rd Qu.:0.0000
## Max.    :1.0000
```

split the dataset into training and testing set

```
library(caTools)
set.seed(209123)

train <- sample.split(NHANES.data$RIDAGEYR, SplitRatio = 0.7)
NHANES.training <- subset(NHANES.data,train==TRUE)
NHANES.testing <- subset(NHANES.data,train==FALSE)

testing.mortstat <- NHANES.testing$mortstat
```

Model1: logistic regression

```
glm.fits <- glm(mortstat ~ ., NHANES.training, family=binomial)
glm.probs <- predict(glm.fits, NHANES.testing, type = "response")
glm.pred <- rep(0,410)
glm.pred[glm.probs > .5] <- 1
```

```
summary(glm.fits)
```

```
##
## Call:
## glm(formula = mortstat ~ ., family = binomial, data = NHANES.training)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.7588  -0.6187  -0.3570  -0.1779   2.5864
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -4.4840607  2.5485249  -1.759  0.07850 .
## RIDAGEYR      0.0826033  0.0140330   5.886 3.95e-09 ***
## RIAGENDR     -0.5584385  0.2131970  -2.619  0.00881 **
## BPQ010        0.0666546  0.1578904   0.422  0.67291
## BPQ060       -0.3399005  0.2699451  -1.259  0.20798
## DIQ010       -0.3104766  0.2279732  -1.362  0.17323
## DIQ050       -0.4984604  0.4032705  -1.236  0.21644
## DIQ090       -0.4797651  0.4372214  -1.097  0.27251
## MCQ010        0.3813123  0.3433004   1.111  0.26669
## MCQ053        0.1093301  0.4560970   0.240  0.81056
## MCQ160A      -0.2799572  0.2037979  -1.374  0.16953
## MCQ160B      -0.6180269  0.3094232  -1.997  0.04579 *
## MCQ160K      -0.1882452  0.3461293  -0.544  0.58654
## MCQ160L      -0.3227434  0.4628207  -0.697  0.48559
## BMXWAIST      0.0002118  0.0003111   0.681  0.49610
## MCQ160M       0.2871458  0.2893780   0.992  0.32106
## MCQ220       -0.1515776  0.2455244  -0.617  0.53700
## MCQ245A       0.2517161  0.2449428   1.028  0.30411
```

```
## MCQ250A      0.1605086  0.1906151   0.842  0.39976
## MCQ250B     -0.1213568  0.2533617  -0.479  0.63195
## MCQ250C     -0.2622214  0.2344844  -1.118  0.26344
## MCQ250E     -0.0956645  0.2484308  -0.385  0.70018
## MCQ250F      0.2430167  0.2156945   1.127  0.25988
## MCQ250G     -0.4689425  0.2637078  -1.778  0.07536 .
## MCQ265      -0.2208634  0.2037583  -1.084  0.27839
## SSQ011      -0.3263575  0.3006307  -1.086  0.27767
## SSQ051      -0.0063933  0.1638842  -0.039  0.96888
## WHQ030       0.2124070  0.1517947   1.399  0.16172
## WHQ040      -0.0547762  0.2356038  -0.232  0.81616
## LBXRDW       0.0347990  0.0077760   4.475  7.63e-06 ***
## HSD010       0.4345622  0.1045889   4.155  3.25e-05 ***
## BPXPULS      0.1183853  0.2976498   0.398  0.69083
## BPXML1       0.0642645  0.0635249   1.012  0.31171
## VIQ200       0.2628171  0.2490539   1.055  0.29131
## BMXBMI       0.0003109  0.0002749   1.131  0.25808
## BPXSY1      -0.0050710  0.0109207  -0.464  0.64240
## BPXDI1      -0.0035235  0.0118823  -0.297  0.76682
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 932.39  on 959  degrees of freedom
## Residual deviance: 729.25  on 923  degrees of freedom
## AIC: 803.25
##
## Number of Fisher Scoring iterations: 5
```

According to logistic regression, RIDAGEYR, RIAGENDR, MCQ160B, LBXRDW, and BPXPULS significantly affect the results.

```
# test for sensitivity and specificity
conf.logistic <- table(glm.pred,testing.mortstat)
conf.logistic

##           testing.mortstat
## glm.pred    0    1
##           0 337  55
##           1   9   9

sens.logistic <- conf.logistic[1,1]/(conf.logistic[1,1]+conf.logistic[1,2])
sens.logistic

## [1] 0.8596939

spec.logistic <- conf.logistic[2,2]/(conf.logistic[2,1]+conf.logistic[2,2])
spec.logistic

## [1] 0.5

# testing error rate
err.logistic <- mean(glm.pred!=testing.mortstat)
err.logistic

## [1] 0.1560976
```

The sensitivity of logistic regression model is 85.97%. The specificity of logistic regression model is 50.00%. The testing error rate of logistic regression model is 15.61%

## Model2: LDA

```
library(MASS)

##
## Attaching package: 'MASS'
## The following object is masked from 'package:dplyr':
##
##      select

lda.fits <- lda(mortstat ~ ., data = NHANES.training)
lda.pred <- predict(lda.fits, NHANES.testing)
lda.class <- lda.pred$class

# test for sensitivity and specificity
conf.lda <- table(lda.class, testing.mortstat)
conf.lda

##           testing.mortstat
## lda.class    0    1
##           0 340  56
##           1   6   8

sens.lda <- conf.lda[1,1]/(conf.lda[1,1]+conf.lda[1,2])
sens.lda

## [1] 0.8585859

spec.lda <- conf.lda[2,2]/(conf.lda[2,1]+conf.lda[2,2])
spec.lda

## [1] 0.5714286

# testing error rate
testing.err.lda <- mean(testing.mortstat!=lda.class)
testing.err.lda

## [1] 0.1512195
```

The sensitivity of LDA model is 85.86%. The specificity of LDA model is 57.14%. The testing error rate of LDA model is 15.12%

## Model3: QDA

```
qda.fits <- qda(mortstat ~ ., data = NHANES.training)
qda.pred <- predict(qda.fits, NHANES.testing)
qda.class <- qda.pred$class

# test for sensitivity and specificity
conf.qda <- table(qda.class, testing.mortstat)
conf.qda

##           testing.mortstat
## qda.class    0    1
##           0 311  46
```

```
##          1  35  18
sens.qda <- conf.qda[1,1]/(conf.qda[1,1]+conf.qda[1,2])
sens.qda
```

```
## [1] 0.8711485
spec.qda <- conf.qda[2,2]/(conf.qda[2,1]+conf.qda[2,2])
spec.qda
```

```
## [1] 0.3396226
# testing error rate
testing.err.qda <- mean(testing.mortstat!=qda.class)
testing.err.qda
```

```
## [1] 0.197561
```

The sensitivity of QDA model is 87.11%. The specificity of QDA model is 33.96%. The testing error rate of QDA model is 19.76%

## Model4: Lasso

```
library(glmnet)
```

```
## Loading required package: Matrix
##
## Attaching package: 'Matrix'
## The following objects are masked from 'package:tidyr':
##
##      expand, pack, unpack
## Loaded glmnet 4.1-3
x.train <- model.matrix(mortstat ~., data = NHANES.training)[,-1]
y.train <- NHANES.training$mortstat

x.test <- model.matrix(mortstat ~ ., data = NHANES.testing)[,-1]
y.test <- NHANES.testing$mortstat

#grid <- 10~seq(10,-2,length=100)

cv.model <- cv.glmnet(x.train,y.train,alpha=1)
best.lambda.lasso <- cv.model$lambda.min
```

```
lasso.mod <- glmnet(x.train, y.train,alpha=1,lambda=best.lambda.lasso)
lasso.val <- as.numeric(predict(lasso.mod, x.test,s=best.lambda.lasso,type="class"))
lasso.pred <- rep(0,410)
lasso.pred[lasso.val > .5] <- 1
```

```
# test for sensitivity and specificity
conf.lasso <- table(lasso.pred,y.test)
conf.lasso
```

```
##          y.test
## lasso.pred  0   1
##           0 345  62
##           1   1   2
```

```
sens.lasso <- conf.lasso[1,1]/(conf.lasso[1,1]+conf.lasso[1,2])
sens.lasso
```

```
## [1] 0.8476658
```

```
spec.lasso <- conf.lasso[2,2]/(conf.lasso[2,1]+conf.lasso[2,2])
spec.lasso
```

```
## [1] 0.6666667
```

```
# MSE
err.lasso <- mean(lasso.pred!=y.test)
err.lasso
```

```
## [1] 0.1536585
```

The sensitivity of QDA model is 84.77%. The specificity of QDA model is 66.67%. The testing error rate of QDA model is 15.37%

## Model 5-7: SVM

```
library(e1071)
set.seed(283)

# linear kernel
tune.out.linear <- tune(svm, mortstat~., data=NHANES.training, kernel="linear", ranges = list(cost=c(0.001, 0.01, 0.1, 1, 10, 100, 1000)))
summary(tune.out.linear)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   1000
##
## - best performance: 0.1743665
##
## - Detailed performance results:
##   cost      error dispersion
## 1 1e-03 0.1761247 0.02178138
## 2 1e-02 0.1758409 0.02225474
## 3 1e-01 0.1757542 0.02238576
## 4 1e+00 0.1757910 0.02241351
## 5 1e+01 0.1757957 0.02240054
## 6 1e+02 0.1756740 0.02248178
## 7 1e+03 0.1743665 0.02212332
tune.out.linear$best.model

##
## Call:
## best.tune(method = svm, train.x = mortstat ~ ., data = NHANES.training,
##   ranges = list(cost = c(0.001, 0.01, 0.1, 1, 10, 100, 1000)),
##   kernel = "linear")
```

```

##
##
## Parameters:
##   SVM-Type:  eps-regression
##   SVM-Kernel: linear
##       cost: 1000
##       gamma: 0.02777778
##       epsilon: 0.1
##
##
## Number of Support Vectors: 499
# svm, linear, cost = 1000
svm.linear <- svm(mortstat ~., data=NHANES.training, kernel ="linear",cost=1000,type="C-classification")
svm.linear.pred <- predict(svm.linear, NHANES.testing)

# test for sensitivity and specificity
conf.svm.lin <- table(svm.linear.pred,testing.mortstat)
conf.svm.lin

##           testing.mortstat
## svm.linear.pred    0     1
##                0 344   62
##                1   2    2

sens.svm.lin <- conf.svm.lin[1,1]/(conf.svm.lin[1,1]+conf.svm.lin[1,2])
sens.svm.lin

## [1] 0.8472906

spec.svm.lin <- conf.svm.lin[2,2]/(conf.svm.lin[2,1]+conf.svm.lin[2,2])
spec.svm.lin

## [1] 0.5

# testing error rate
err.svm.lin <- mean((svm.linear.pred!=testing.mortstat))
err.svm.lin

## [1] 0.1560976

The sensitivity of SVM model using linear kernel is 84.73%. The specificity of SVM model using linear kernel
is 50.00%. The testing error of SVM model using linear kernel is 15.61%

# radial kernel
set.seed(32989)
tune.out.radial <- tune(svm, mortstat~., data=NHANES.training, kernel ="radial", ranges =list(cost=c(0.1,1000),gamma=c(0.001,0.1),epsilon=c(0.001,0.1)))
summary(tune.out.radial)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost gamma
##    10     3
##

```



```
## - best performance: 0.1546424
##
## - Detailed performance results:
##      cost gamma      error dispersion
## 1  1e-03    0.5 0.1762250 0.04128445
## 2  1e-02    0.5 0.1760117 0.04123115
## 3  1e-01    0.5 0.1736311 0.04058846
## 4  1e+00    0.5 0.1574270 0.03356227
## 5  1e+01    0.5 0.1546898 0.02555857
## 6  1e+02    0.5 0.1546898 0.02555857
## 7  1e+03    0.5 0.1546898 0.02555857
## 8  1e-03    1.0 0.1762251 0.04128443
## 9  1e-02    1.0 0.1760122 0.04123087
## 10 1e-01    1.0 0.1736418 0.04058795
## 11 1e+00    1.0 0.1574586 0.03353923
## 12 1e+01    1.0 0.1546960 0.02543562
## 13 1e+02    1.0 0.1546960 0.02543562
## 14 1e+03    1.0 0.1546960 0.02543562
## 15 1e-03    2.0 0.1762251 0.04128453
## 16 1e-02    2.0 0.1760119 0.04123084
## 17 1e-01    2.0 0.1736399 0.04058787
## 18 1e+00    2.0 0.1574365 0.03352892
## 19 1e+01    2.0 0.1546469 0.02540100
## 20 1e+02    2.0 0.1546469 0.02540100
## 21 1e+03    2.0 0.1546469 0.02540100
## 22 1e-03    3.0 0.1762251 0.04128452
## 23 1e-02    3.0 0.1760119 0.04123084
## 24 1e-01    3.0 0.1736396 0.04058779
## 25 1e+00    3.0 0.1574341 0.03352662
## 26 1e+01    3.0 0.1546424 0.02539442
## 27 1e+02    3.0 0.1546424 0.02539442
## 28 1e+03    3.0 0.1546424 0.02539442
## 29 1e-03    4.0 0.1762251 0.04128452
## 30 1e-02    4.0 0.1760119 0.04123084
## 31 1e-01    4.0 0.1736396 0.04058777
## 32 1e+00    4.0 0.1574341 0.03352617
## 33 1e+01    4.0 0.1546424 0.02539304
## 34 1e+02    4.0 0.1546424 0.02539304
## 35 1e+03    4.0 0.1546424 0.02539304
```

```
tune.out.radial$best.model
```

```
##
## Call:
## best.tune(method = svm, train.x = mortstat ~ ., data = NHANES.training,
##      ranges = list(cost = c(0.001, 0.01, 0.1, 1, 10, 100, 1000), gamma = c(0.5,
##      1, 2, 3, 4)), kernel = "radial")
##
##
## Parameters:
##      SVM-Type:  eps-regression
##      SVM-Kernel: radial
##      cost:      10
##      gamma:     3
##      epsilon:   0.1
```

```
##
##
## Number of Support Vectors: 960
set.seed(427)

# svm, radial, cost = 10, gamma=3
svm.radial <- svm(mortstat ~., data=NHANES.training, kernel ="radial",cost=10,gamma=3,type="C-classification")
svm.radial.pred <- predict(svm.radial, NHANES.testing)

# test for sensitivity and specificity
conf.svm.rad <- table(svm.radial.pred,testing.mortstat)
conf.svm.rad
```

```
##           testing.mortstat
## svm.radial.pred  0    1
##                0 346  64
##                1   0   0

sens.svm.rad <- conf.svm.rad[1,1]/(conf.svm.rad[1,1]+conf.svm.rad[1,2])
sens.svm.rad
```

```
## [1] 0.8439024
```

```
spec.svm.rad <- conf.svm.rad[2,2]/(conf.svm.rad[2,1]+conf.svm.rad[2,2])
spec.svm.rad
```

```
## [1] NaN
```

```
# testing error rate
err.svm.rad <- mean((svm.radial.pred!=testing.mortstat))
err.svm.rad
```

```
## [1] 0.1560976
```

The sensitivity of SVM model using radial kernel is 84.39%. The specificity of SVM model using radial kernel is unknown since the model classify all cases as alive. The testing error of SVM model using radial kernel is 15.61%

```
# polynomial kernel
set.seed(2987)
tune.out.poly <- tune(svm, mortstat~., data=NHANES.training, , kernel ="polynomial", ranges = list(cost = 1e-03, degree = 1:3))
summary(tune.out.poly)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost degree
##     1     3
##
## - best performance: 0.1605922
##
## - Detailed performance results:
##   cost degree   error dispersion
## 1 1e-03      1 0.1762409 0.03137452
```

```
## 2 1e-02      1 0.1761873 0.03141165
## 3 1e-01      1 0.1758970 0.03179227
## 4 1e+00      1 0.1754025 0.03251923
## 5 5e+00      1 0.1753910 0.03248814
## 6 1e+01      1 0.1753903 0.03250287
## 7 1e+02      1 0.1753745 0.03249319
## 8 1e-03      2 0.1761841 0.03136474
## 9 1e-02      2 0.1754535 0.03128604
## 10 1e-01     2 0.1694481 0.03049072
## 11 1e+00     2 0.1704060 0.02605340
## 12 5e+00     2 0.2321557 0.03014433
## 13 1e+01     2 0.2887254 0.04164993
## 14 1e+02     2 0.7398740 0.16085113
## 15 1e-03     3 0.1761005 0.03134433
## 16 1e-02     3 0.1745599 0.03106061
## 17 1e-01     3 0.1639244 0.02911744
## 18 1e+00     3 0.1605922 0.02318568
## 19 5e+00     3 0.1821862 0.02972759
## 20 1e+01     3 0.1996163 0.03122904
## 21 1e+02     3 0.2721005 0.06066044
## 22 1e-03     4 0.1760316 0.03132022
## 23 1e-02     4 0.1739593 0.03087355
## 24 1e-01     4 0.1633753 0.02895433
## 25 1e+00     4 0.1623101 0.02073873
## 26 5e+00     4 0.1698467 0.02643086
## 27 1e+01     4 0.1790407 0.03106939
## 28 1e+02     4 0.2106294 0.03532869
## 29 1e-03     5 0.1759473 0.03130311
## 30 1e-02     5 0.1733903 0.03080272
## 31 1e-01     5 0.1647294 0.02928668
## 32 1e+00     5 0.1631628 0.02242649
## 33 5e+00     5 0.1659554 0.02470952
## 34 1e+01     5 0.1735138 0.03044052
## 35 1e+02     5 0.2184457 0.08200631
```

```
tune.out.poly$best.model
```

```
##
## Call:
## best.tune(method = svm, train.x = mortstat ~ ., data = NHANES.training,
##   ranges = list(cost = c(0.001, 0.01, 0.1, 1, 5, 10, 100), degree = c(1,
##   2, 3, 4, 5)), kernel = "polynomial")
##
##
## Parameters:
##   SVM-Type:  eps-regression
##   SVM-Kernel: polynomial
##     cost:    1
##   degree:    3
##   gamma:     0.02777778
##   coef.0:    0
##   epsilon:   0.1
##
##
## Number of Support Vectors: 680
```

```

# svm, polynomial, cost = 1, degree=3
svm.poly <- svm(mortstat ~., data=NHANES.training, kernel ="polynomial",cost=1,degree=3,type="C-classif
svm.poly.pred <- predict(svm.poly, NHANES.testing)

# test for sensitivity and specificity
conf.svm.poly <- table(svm.poly.pred,testing.mortstat)
conf.svm.poly

##           testing.mortstat
## svm.poly.pred    0    1
##           0 338   60
##           1   8    4

sens.svm.poly <- conf.svm.poly[1,1]/(conf.svm.poly[1,1]+conf.svm.poly[1,2])
sens.svm.poly

## [1] 0.8492462

spec.svm.poly <- conf.svm.poly[2,2]/(conf.svm.poly[2,1]+conf.svm.poly[2,2])
spec.svm.poly

## [1] 0.3333333
# testing error rate
err.svm.poly <- mean((svm.poly.pred!=testing.mortstat))
err.svm.poly

## [1] 0.1658537

```

The sensitivity of SVM model using polynomial kernel is 84.92%. The specificity of SVM model using polynomial kernel is 33.33%. The testing error of SVM model using polynomial kernel is 16.59%

## Model 8&9: Tree-based methods

### bagging

```

library(tree)
library(randomForest)

## randomForest 4.6-14
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
## The following object is masked from 'package:dplyr':
##
##      combine
## The following object is masked from 'package:ggplot2':
##
##      margin
set.seed(7321)

bag.mod <- randomForest(mortstat~., NHANES.training, mtry=36, ntree=100)

## Warning in randomForest.default(m, y, ...): The response has five or fewer

```

```
## unique values. Are you sure you want to do regression?
```

```
yhat.bag <- predict(bag.mod, NHANES.testing)
```

```
yhat.pred <- rep(0,410)
yhat.pred[yhat.bag > .5] <- 1
```

```
# test for sensitivity and specificity
conf.bag <- table(yhat.pred,testing.mortstat)
conf.bag
```

```
##           testing.mortstat
## yhat.pred    0    1
##           0 333  47
##           1  13  17
```

```
sens.bag <- conf.bag[1,1]/(conf.bag[1,1]+conf.bag[1,2])
sens.bag
```

```
## [1] 0.8763158
```

```
spec.bag <- conf.bag[2,2]/(conf.bag[2,1]+conf.bag[2,2])
spec.bag
```

```
## [1] 0.5666667
```

```
# testing error rate
err.bag <- mean(yhat.pred!=testing.mortstat)
err.bag
```

```
## [1] 0.1463415
```

The sensitivity using bagging is 87.63%. The specificity using bagging is 56.67%. The testing error using bagging is 14.63%.

## boosting

```
library(gbm)
```

```
## Loaded gbm 2.1.8
```


```
set.seed(3253)
```

```
a <- seq(-10, -1, by=0.5)
lambdas <- 10^a
MSE.training <- c()
MSE.testing <- c()
```

```
for (i in 1:length(lambdas)){
  boosting.mod <- gbm(mortstat~., data = NHANES.training, distribution = "gaussian", n.trees = 1000, inter
  boosting.training.pred <- predict(boosting.mod, NHANES.training,n.trees=1000)
  boosting.testing.pred <- predict(boosting.mod, NHANES.testing,n.trees=1000)

  MSE.training[i] <- mean((NHANES.training$mortstat-boosting.training.pred)^2)
  MSE.testing[i] <- mean((testing.mortstat-boosting.testing.pred)^2)
}
```

```
boosting.best <- gbm(mortstat~., data = NHANES.training, distribution = "gaussian", n.trees = 1000, int
summary(boosting.best)
```



| ## | var      | rel.inf              |
|----|----------|----------------------|
| ## | RIDAGEYR | RIDAGEYR 26.88195989 |
| ## | LBXRDW   | LBXRDW 18.12641161   |
| ## | HSD010   | HSD010 9.22796158    |
| ## | BMXBMI   | BMXBMI 8.62360862    |
| ## | BPXDI1   | BPXDI1 7.04995409    |
| ## | BMXWAIST | BMXWAIST 4.89434807  |
| ## | BPXSY1   | BPXSY1 4.46338145    |
| ## | MCQ160B  | MCQ160B 3.26020083   |
| ## | RIAGENDR | RIAGENDR 2.72896902  |
| ## | BPXML1   | BPXML1 2.44635559    |
| ## | DIQ010   | DIQ010 2.21414646    |
| ## | DIQ050   | DIQ050 1.30179636    |
| ## | MCQ245A  | MCQ245A 0.89538498   |
| ## | MCQ160A  | MCQ160A 0.87741648   |
| ## | MCQ250A  | MCQ250A 0.87220896   |
| ## | MCQ220   | MCQ220 0.78822142    |
| ## | WHQ030   | WHQ030 0.71426773    |
| ## | DIQ090   | DIQ090 0.68618030    |
| ## | SSQ051   | SSQ051 0.48661673    |
| ## | WHQ040   | WHQ040 0.43280434    |
| ## | BPXPULS  | BPXPULS 0.42937757   |
| ## | MCQ250F  | MCQ250F 0.41754843   |
| ## | VIQ200   | VIQ200 0.41624430    |
| ## | MCQ265   | MCQ265 0.29331061    |
| ## | MCQ250G  | MCQ250G 0.23668833   |
| ## | MCQ250E  | MCQ250E 0.21502059   |
| ## | MCQ250C  | MCQ250C 0.19787488   |
| ## | MCQ250B  | MCQ250B 0.15638291   |

```
## MCQ160M    MCQ160M    0.14331229
## MCQ010     MCQ010     0.12635022
## BPQ060     BPQ060     0.11641954
## MCQ160K    MCQ160K    0.10094283
## BPQ010     BPQ010     0.05856488
## MCQ053     MCQ053     0.04124174
## MCQ160L    MCQ160L    0.04063063
## SSQ011     SSQ011     0.03789570

boosting.probs <- predict(boosting.best, NHANES.testing, n.trees=1000)
boosting.pred <- rep(0,410)
boosting.pred[boosting.probs > .5] <- 1
```

```
# test for sensitivity and specificity
conf.boost <- table(boosting.pred,testing.mortstat)
conf.boost
```

```
##           testing.mortstat
## boosting.pred  0    1
##              0 342  51
##              1   4  13

sens.boost <- conf.boost[1,1]/(conf.boost[1,1]+conf.boost[1,2])
sens.boost
```

```
## [1] 0.870229

spec.boost <- conf.boost[2,2]/(conf.boost[2,1]+conf.boost[2,2])
spec.boost
```

```
## [1] 0.7647059

# testing error rate
err.boost <- mean(boosting.pred!=testing.mortstat)
err.boost
```

```
## [1] 0.1341463
```

The sensitivity using boosting is 87.02%. The specificity using boosting is 76.47%. The testing error using boosting is 13.41%.

## Comparing model evaluations

```
a <- rbind(c(err.logistic, testing.err.lda, testing.err.qda, err.lasso, err.svm.lin, err.svm.rad, err.svm.pois,
             c(sens.logistic, sens.lda, sens.qda, sens.lasso, sens.svm.lin, sens.svm.rad, sens.svm.pois),
             c(spec.logistic, spec.lda, spec.qda, spec.lasso, spec.svm.lin, spec.svm.rad, spec.svm.pois)),
           print(a)
```

```
##           [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
## [1,] 0.1560976 0.1512195 0.1975610 0.1536585 0.1560976 0.1560976 0.1658537
## [2,] 0.8596939 0.8585859 0.8711485 0.8476658 0.8472906 0.8439024 0.8492462
## [3,] 0.5000000 0.5714286 0.3396226 0.6666667 0.5000000      NaN 0.3333333
##           [,8]      [,9]
## [1,] 0.1463415 0.1341463
## [2,] 0.8763158 0.8702290
## [3,] 0.5666667 0.7647059
```

The results show testing error rate, sensitivity, and specificity for the 9 models respectively, using the list of 36 predictors. Model 8 (bagging) has the highest sensitivity, model 9 (boosting) has the highest specificity as

well as the lowest error rate (highest accuracy). Overall, boosting method is optimal for classifying mortality in the NHANES dataset.

The boosting model achieves sensitivity of 87.02% and specificity of 76.47%, with a testing error rate of 13.41%. Using this model, the top 5 important variables in prediction are: “RIDAGEYR”, “LBXRDW”, “HSD010”, “BMXBMI”, and “BPXDI1”.